



Regularizing Neural Networks via Adversarial Model Perturbation

Yaowei Zheng,¹ Richong Zhang,¹ Yongyi Mao²

¹BDBC and SKLSDE, Beihang University, Beijing, China

²School of EECS, University of Ottawa, Ottawa, Canada




Summary

Background: Effective regularization schemes are highly desired in deep learning for alleviating overfitting and improving generalization.

Motivation: Previous work suggested that flat minima can improve generalization both in theoretical and empirical perspectives [1-4].

Contribution: We propose Adversarial Model Perturbation (AMP) as a powerful regularization scheme principled by the objective of finding flat minima, which achieves better classification performance and calibration results.



(Better minima are flatter, visualized by Li *et al.*, 2018.)

Training Algorithm

Optimization Objective:

$$\min_{\theta} \max_{\Delta: \|\Delta\| \leq \epsilon} \mathcal{L}_{\text{ERM}}(\theta + \Delta)$$

Pseudo-code:

```

while  $\theta$  not converged do
  Initialize perturbation  $\Delta$  with 0;
  for  $n \leftarrow 1$  to  $N$  do
    Update  $\Delta$  to maximize  $\mathcal{L}_{\text{ERM}}(\theta + \Delta)$  via
    gradient ascent with learning rate  $\zeta$ ;
    if  $\|\Delta\|_2 > \epsilon$  then
      Normalize  $\Delta$  to restrict its norm  $\|\Delta\|_2$  to  $\epsilon$ ;
  Update  $\theta$  to minimize  $\mathcal{L}_{\text{ERM}}(\theta + \Delta)$  via gradient
  descent with learning rate  $\eta$ ;

```

PyTorch Implementation:

```

from amp import AMP
opt = AMP(params, lr=0.1, epsilon=0.5)
for inputs, targets in dataset:
  def closure():
    opt.zero_grad()
    outputs = model(inputs)
    loss = loss_fn(outputs, targets)
    loss.backward()
    return outputs, loss
  outputs, loss = opt.step(closure)

```

Official Code:

<https://github.com/hiyouga/AMP-Regularizer>

Method


AMP: Adversarial Model Perturbation

We derive an AMP loss from the empirical risk (ERM loss) by applying the “worst” perturbation on the model parameters to penalize the sharp minima.

$$\mathcal{L}_{\text{ERM}}(\theta) := \frac{1}{|D|} \sum_{(x,y) \in D} \ell(x, y; \theta)$$

$$\mathcal{L}_{\text{AMP}}(\theta) := \max_{\Delta: \|\Delta\| \leq \epsilon} \mathcal{L}_{\text{ERM}}(\theta + \Delta)$$

It can be seen as an analogue of a “max-pooling” operation on the empirical risk on each point in the parameter space.



Theoretical Justification


AMP Finds Flatter Minima:

We assume that the loss surface of each local minimum in \mathcal{L}_{ERM} can be *locally* approximated as an inverted Gaussian surface γ with a mean vector μ and a covariance matrix κ . Under the locally Gaussian assumption, the empirical risk $\gamma(\theta; \mu, \kappa, A, C)$ is minimized when $\theta = \mu$ and the minimum value is $\gamma^*(\mu, \kappa, A, C) = C - A$. The minimum value of the AMP loss is the empirical risk at the location in the narrowest principal direction of the cross-section of the loss surface:

$$\gamma_{\text{AMP}}^*(\mu, \kappa, A, C) = C - A \exp\left(-\frac{\epsilon^2}{2\sigma^2}\right)$$

where σ^2 is the smallest eigenvalue of κ .

It is clear that γ_{AMP}^* although related to the minimum value of empirical risk γ^* , it also takes into account the curvature of the surface around the local minimum.



AMP Regularizes Gradient Norm:

Consider that $N = 1$ (in fact used). The AMP training is equivalent to ERM training with an additional term:

$$\tilde{\mathcal{J}}_{\text{ERM}}(\theta) := \mathcal{J}_{\text{ERM}}(\theta) + \Omega(\theta)$$

where

$$\Omega(\theta) := \begin{cases} \zeta \|\nabla_{\theta} \mathcal{J}_{\text{ERM}}(\theta)\|_2^2, & \|\zeta \nabla_{\theta} \mathcal{J}_{\text{ERM}}(\theta)\|_2 \leq \epsilon \\ \epsilon \|\nabla_{\theta} \mathcal{J}_{\text{ERM}}(\theta)\|_2, & \|\zeta \nabla_{\theta} \mathcal{J}_{\text{ERM}}(\theta)\|_2 > \epsilon \end{cases}$$

Note that a minimum with smaller gradient norms around it is a flatter minimum.

Experiments

Results on Image Classification Benchmarks:

PreActResNet18	Test Error (%)	Test NLL
ERM	2.95±0.063	0.166±0.004
Dropout	2.80±0.065	0.156±0.012
Label Smoothing	2.78±0.087	0.998±0.002
Flooding	2.84±0.047	0.130±0.003
MixUp	2.74±0.044	0.146±0.004
Adv. Training	2.77±0.080	0.151±0.018
RMP	2.93±0.066	0.161±0.010
AMP	2.30±0.025	0.096±0.002

(a) SVHN

PreActResNet18	Test Error (%)	Test NLL
ERM	5.02±0.212	0.239±0.009
Dropout	4.86±0.148	0.223±0.009
Label Smoothing	4.85±0.115	1.038±0.003
Flooding	4.97±0.082	0.166±0.003
MixUp	4.09±0.117	0.198±0.004
Adv. Training	4.99±0.085	0.247±0.006
RMP	4.97±0.167	0.239±0.008
AMP	3.97±0.091	0.129±0.003

(b) CIFAR-10


PreActResNet18	Test Error (%)	Test NLL
ERM	24.31±0.303	1.056±0.013
Dropout	24.48±0.351	1.110±0.021
Label Smoothing	22.07±0.256	2.099±0.005
Flooding	24.50±0.234	0.950±0.011
MixUp	21.78±0.210	0.910±0.007
Adv. Training	25.23±0.229	1.110±0.012
RMP	24.28±0.138	1.059±0.011
AMP	21.51±0.308	0.774±0.016

(c) CIFAR-100


Improvement over Various Data Augmentation Techniques:

Test Error (%)	Augment	WideResNet-28-10		PyramidNet-164-270	
Dataset	Technique	ERM	AMP	ERM	AMP
SVHN	Vanilla	2.57±0.067	2.19±0.036	2.47±0.034	2.11±0.041
	Cutout	2.27±0.085	1.83±0.018	2.19±0.021	1.82±0.023
	AutoAug	1.91±0.059	1.61±0.024	1.80±0.044	1.35±0.056
CIFAR-10	Vanilla	3.87±0.167	3.00±0.059	3.60±0.197	2.75±0.040
	Cutout	3.38±0.081	2.67±0.043	2.83±0.102	2.27±0.034
	AutoAug	2.78±0.134	2.32±0.097	2.49±0.128	1.98±0.062
CIFAR-100	Vanilla	19.17±0.270	17.33±0.110	17.13±0.210	15.09±0.092
	Cutout	18.12±0.114	16.04±0.071	16.45±0.136	14.34±0.153
	AutoAug	17.79±0.185	14.95±0.088	15.43±0.269	13.36±0.245


Calibration Results: (lower is better)




Flatness of the Selected Minima:




ERM Training Loss



ERM Test Loss




AMP Training Loss




AMP Test Loss

Loss Values with Varying Perturbation Size:



(a) CIFAR-10 Training Set



(b) CIFAR-10 Test Set

References:

- [1] Hochreiter *et al.* Flat minima. Neural Computation, 9(1):1-42, 1997.
- [2] Keskar *et al.* On large-batch training for deep learning: Generalization gap and sharp minima. In ICLR, 2017.
- [3] Li *et al.* Visualizing the loss landscape of neural nets. In NeurIPS, 2018.
- [4] Foret *et al.* Sharpness-aware minimization for efficiently improving generalization. In ICLR, 2021.