



Regularizing Neural Networks via Adversarial Model Perturbation

Yaowei Zheng¹ Richong Zhang¹ Yongyi Mao²

¹BDBC and SKLSDE, Beihang University, Beijing, China


²School of EECS, University of Ottawa, Ottawa, Canada

CVPR 2021

- 1 Background
- 2 AMP: Adversarial Model Perturbation
- 3 Theoretical Justifications of AMP
- 4 Experiments
- 5 Conclusion

Regularization Alleviate Overfitting


Effective regularization schemes alleviate overfitting and improve generalization.




- Some researchers have found the modern neural networks may have different behaviour, *i.e.*, the *Double Descent* (Nakkiran *et al.*, 2020). Nevertheless, well-regularized neural networks consistently achieve better performance in practice.

Flat Minima Helps Generalization

- Flat minima correspond to low-complexity networks. (Hochreiter *et al.*, 1997)
- Small-batch SGD produces flat minima that generalize well. (Keskar *et al.*, 2017)
- Better minimizers of loss function are flatter in visualization. (Li *et al.*, 2018)
- A PAC-Bayes based generalization guarantee for flat minima. (Foret *et al.*, 2020)



(a) ResNet without skip connections



(b) ResNet with skip connections (Li *et al.*, 2018)



From Empirical Risk to AMP Loss

The AMP loss is derived from empirical risk by applying the “worst” perturbation on the model parameters.

$$\mathcal{L}_{\text{ERM}}(\boldsymbol{\theta}) := \frac{1}{|\mathcal{D}|} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}} \ell(\mathbf{x}, \mathbf{y}; \boldsymbol{\theta}) \quad (1)$$

$$\mathcal{L}_{\text{AMP}}(\boldsymbol{\theta}) := \max_{\Delta: \|\Delta\| \leq \epsilon} \mathcal{L}_{\text{ERM}}(\boldsymbol{\theta} + \Delta) \quad (2)$$

As sketched in the figures, it applies a “max-pooling” operation on the empirical risk to seek a flatter minimum.



Training Algorithm

A mini-batch SGD is used for solving the “min-max” problem.

$$\min_{\theta} \max_{\Delta: \|\Delta\| \leq \epsilon} \mathcal{L}_{\text{ERM}}(\theta + \Delta) \quad (3)$$

Algorithm 1: Adversarial Model Perturbation Training

```
1 while  $\theta$  not converged do
2   Initialize perturbation  $\Delta$  with 0;
3   for  $n \leftarrow 1$  to  $N$  do
4     Update  $\Delta$  to maximize  $\mathcal{L}_{\text{ERM}}(\theta + \Delta)$  via gradient ascent with learning rate  $\zeta$ ;
5     if  $\|\Delta\|_2 > \epsilon$  then
6       Normalize  $\Delta$  to restrict its norm  $\|\Delta\|_2$  to  $\epsilon$ ;
7   Update  $\theta$  to minimize  $\mathcal{L}_{\text{ERM}}(\theta + \Delta)$  via gradient descent with learning rate  $\eta$ ;
```

Implementation

Source code: <https://github.com/hiyouga/AMP-Regularizer>

```
1 from amp import AMP
2 optimizer = AMP(model.parameters(), lr=0.1, epsilon=0.5, momentum=0.9)
3 for inputs, targets in dataset:
4     def closure():
5         optimizer.zero_grad()
6         outputs = model(inputs)
7         loss = loss_fn(outputs, targets)
8         loss.backward()
9         return outputs, loss
10    outputs, loss = optimizer.step(closure)
```

AMP Finds Flatter Local Minima

We assume that the loss surface of each local minimum in \mathcal{L}_{ERM} can be *locally* approximated as an inverted Gaussian surface γ with a mean vector μ and a covariance matrix κ .

Theorem (informal)

Under the locally Gaussian assumption, the empirical risk $\gamma(\theta; \mu, \kappa, A, C)$ is minimized when $\theta = \mu$ and the minimum value is $\gamma^(\mu, \kappa, A, C) = C - A$. The minimum value of the AMP loss is the empirical risk at the location in the narrowest principal direction of the cross-section of the loss surface:*

$$\gamma_{\text{AMP}}^*(\mu, \kappa, A, C) = C - A \exp\left(-\frac{\epsilon^2}{2\sigma^2}\right) \quad (4)$$

where σ^2 is the smallest eigenvalue of κ .





Figure: The minimum values of γ and γ_{AMP} .

AMP Regularizes Gradient Norm

Theorem (informal)

Consider that $N = 1$, which is in fact used in our experiments. The AMP training is equivalent to ERM training with an additional term:

$$\tilde{\mathcal{J}}_{\text{ERM}}(\boldsymbol{\theta}) := \mathcal{J}_{\text{ERM}}(\boldsymbol{\theta}) + \Omega(\boldsymbol{\theta}) \quad (5)$$

where

$$\Omega(\boldsymbol{\theta}) := \begin{cases} \zeta \|\nabla_{\boldsymbol{\theta}} \mathcal{J}_{\text{ERM}}(\boldsymbol{\theta})\|_2^2, & \|\zeta \nabla_{\boldsymbol{\theta}} \mathcal{J}_{\text{ERM}}(\boldsymbol{\theta})\|_2 \leq \epsilon \\ \epsilon \|\nabla_{\boldsymbol{\theta}} \mathcal{J}_{\text{ERM}}(\boldsymbol{\theta})\|_2, & \|\zeta \nabla_{\boldsymbol{\theta}} \mathcal{J}_{\text{ERM}}(\boldsymbol{\theta})\|_2 > \epsilon \end{cases} \quad (6)$$

Experimental Setup

Image classification datasets:

- SVHN (10-way)
- CIFAR-10 (10-way)
- CIFAR-100 (100-way)


Compared methods:

- ERM (Vapnik *et al.*, 1998)
- Dropout (Srivastava *et al.*, 2014)
- Label smoothing (Szegedy *et al.*, 2016)
- Flooding (Ishida *et al.*, 2020)
- MixUp (Zhang *et al.*, 2018)
- Adversarial Training (Goodfellow *et al.*, 2015)


Model architectures:

- PreActResNet18 (He *et al.*, 2016)
- VGG16 (Simonyan *et al.*, 2015)
- WideResNet-28-10 (Zagoruyko *et al.*, 2016)
- PyramidNet-164-270 (Han *et al.*, 2017)

Loss Curves



(a) CIFAR-10 Training Set



(b) CIFAR-10 Test Set

Results on Image Classification Benchmarks

PreActResNet18	Test Error (%)	Test NLL	PreActResNet18	Test Error (%)	Test NLL	PreActResNet18	Test Error (%)	Test NLL
ERM	2.95±0.063	0.166±0.004	ERM	5.02±0.212	0.239±0.009	ERM	24.31±0.303	1.056±0.013
Dropout	2.80±0.065	0.156±0.012	Dropout	4.86±0.148	0.223±0.009	Dropout	24.48±0.351	1.110±0.021
Label Smoothing	2.78±0.087	0.998±0.002	Label Smoothing	4.85±0.115	1.038±0.003	Label Smoothing	22.07±0.256	2.099±0.005
Flooding	2.84±0.047	0.130±0.003	Flooding	4.97±0.082	0.166±0.003	Flooding	24.50±0.234	0.950±0.011
MixUp	2.74±0.044	0.146±0.004	MixUp	4.09±0.117	0.198±0.004	MixUp	21.78±0.210	0.910±0.007
Adv. Training	2.77±0.080	0.151±0.018	Adv. Training	4.99±0.085	0.247±0.006	Adv. Training	25.23±0.229	1.110±0.012
RMP	2.93±0.066	0.161±0.010	RMP	4.97±0.167	0.239±0.008	RMP	24.28±0.138	1.059±0.011
AMP	2.30±0.025	0.096±0.002	AMP	3.97±0.091	0.129±0.003	AMP	21.51±0.308	0.774±0.016
VGG16	Test Error (%)	Test NLL	VGG16	Test Error (%)	Test NLL	VGG16	Test Error (%)	Test NLL
ERM	3.14±0.060	0.140±0.027	ERM	6.32±0.193	0.361±0.012	ERM	27.84±0.297	1.827±0.209
Dropout	2.96±0.049	0.134±0.027	Dropout	6.22±0.147	0.314±0.009	Dropout	27.72±0.337	1.605±0.062
Label Smoothing	3.07±0.070	1.004±0.002	Label Smoothing	6.29±0.158	1.076±0.003	Label Smoothing	27.49±0.179	2.310±0.005
Flooding	3.15±0.085	0.128±0.003	Flooding	6.26±0.145	0.234±0.005	Flooding	27.93±0.271	1.221±0.037
MixUp	3.09±0.057	0.160±0.003	MixUp	5.48±0.112	0.251±0.003	MixUp	<u>26.81±0.254</u>	<u>1.136±0.013</u>
Adv. Training	<u>2.94±0.091</u>	<u>0.122±0.003</u>	Adv. Training	6.49±0.130	0.380±0.010	Adv. Training	29.12±0.145	1.535±0.389
RMP	3.19±0.052	0.134±0.004	RMP	6.30±0.109	0.363±0.010	RMP	27.81±0.327	1.873±0.035
AMP	2.73±0.015	0.116±0.006	AMP	<u>5.65±0.147</u>	0.207±0.005	AMP	25.60±0.168	1.049±0.049

(a) SVHN

(b) CIFAR-10

(c) CIFAR-100

Table: Top-1 classification errors and test neg-log-likelihoods.

Improvement over Data Augmentation


		WideResNet-28-10		PyramidNet-164-270	
		ERM	AMP	ERM	AMP
SVHN	Vanilla	2.57 ± 0.067	2.19 ± 0.036	2.47 ± 0.034	2.11 ± 0.041
	Cutout	2.27 ± 0.085	1.83 ± 0.018	2.19 ± 0.021	1.82 ± 0.023
	AutoAug	1.91 ± 0.059	1.61 ± 0.024	1.80 ± 0.044	1.35 ± 0.056
CIFAR-10	Vanilla	3.87 ± 0.167	3.00 ± 0.059	3.60 ± 0.197	2.75 ± 0.040
	Cutout	3.38 ± 0.081	2.67 ± 0.043	2.83 ± 0.102	2.27 ± 0.034
	AutoAug	2.78 ± 0.134	2.32 ± 0.097	2.49 ± 0.128	1.98 ± 0.062
CIFAR-100	Vanilla	19.17 ± 0.270	17.33 ± 0.110	17.13 ± 0.210	15.09 ± 0.092
	Cutout	18.12 ± 0.114	16.04 ± 0.071	16.45 ± 0.136	14.34 ± 0.153
	AutoAug	17.79 ± 0.185	14.95 ± 0.088	15.43 ± 0.269	13.36 ± 0.245

Table: Top-1 classification errors and test neg-log-likelihoods.

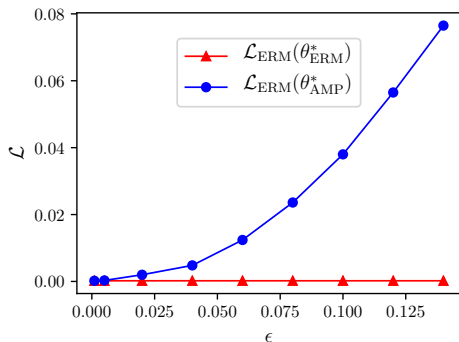
Calibration Results

Expected Calibration Error: (lower is better)


$$\text{ECE} = \sum_{m=1}^M \frac{|B_m|}{n} \left| \underbrace{\frac{1}{|B_m|} \sum_{i \in B_m} 1(\hat{y}_i = y_i)}_{\text{accuracy}} - \underbrace{\frac{1}{|B_m|} \sum_{i \in B_m} \hat{p}_i}_{\text{confidence}} \right|$$



Loss Values with Varying Perturbation Size




(a) CIFAR-10 Training Set




(b) CIFAR-10 Test Set


Flatness of the Selected Minima




(a) ERM Training Loss



(b) ERM Test Loss



(c) AMP Training Loss



(d) AMP Test Loss

- Motivated by the understanding that flat minima help generalization, we propose adversarial model perturbation (AMP) as an efficient regularization scheme.
- We theoretically justify that AMP is capable of finding flatter local minima, thereby improving generalization.
- Extensive experiments on the benchmark datasets demonstrate that AMP achieves the best performance among the compared regularization schemes on various modern neural network architectures.

ArXiv: <https://arxiv.org/abs/2010.04925>

Code: <https://github.com/hiyouga/AMP-Regularizer>